

Optimization Formalization

Sara Achour

May 29, 2014

1 Model V2

1.1 Hardware Model

Worker Machine $m \in M, |M|$
Main machine R

Machine cost for task with runtime t :

$$c = a \cdot W \cdot t + b \cdot E \cdot t$$

E: energy per unit time.

W: amount of work per unit time.

a: weight to apply to machine efficiency.

b: weight to apply to machine power consumption.

p_{fail} : probability of a crash per unit time.

p_{hang} : probability of a hang per unit time.

p_{err} : probability of soft data error occurring per unit time.

$$p_{ok} = 1 - p_{fail} + p_{hang} + p_{err}$$

Oversimplification: Assuming the probability of various failures is not dependent on the computation running.

1.2 Programming Model

taskset T :

size: $|T|$

number inputs: n

number outputs: m

task $t \in T$:

$$t: R^m \rightarrow R^n$$

task t is over real numbers. maps m dimension vector to n dimension vector.

$$t: f(i_1, \cdot, i_n) \rightarrow (r_1, \cdot, r_m)$$

i_k : k th input, let \vec{i} be input vector.

r_k : kth output, let \vec{r} be output vector.
 $P(\vec{r}, \vec{i})$
 The probability distribution of the outputs and inputs.
 $t: f(i_1, \dots, i_n) \sim \operatorname{argmax}_{\hat{r}} P(\hat{r} | \vec{i})$
 The input-output relation f can be rewritten as the expectation
 maximization of r given i .
 Since f is deterministic, there should be a particular r where $P(r, i) = 1$
 for every i .
 ex: $P(\vec{r} | \vec{i})$
 The probability of outputs r occurring, given input i .
 δ_t : the runtime of the particular task.

1.3 Hardware-Programming Model Interaction

For a particular task t running on machine m :

$p_{fail}^t = \delta \cdot p_{fail}$
 $p_{hang}^t = \delta \cdot p_{hang}$
 $p_{err}^t = \delta \cdot p_{err}$
 $P_{err}(\vec{r}, \vec{i})$
 This is the output distribution of the task if an error occurs.
 $P'(\vec{r}, \vec{i}) = P(\vec{r}, \vec{i}) + P_{err}(\vec{r}, \vec{i})$
 the output distribution on machine m (P') is a mixture of the
 output distribution (P) distribution and the error distribution (P_{err})
 .
 $P_\epsilon(\epsilon)$
 this is the distribution of errors between the correct, faulty
 result. Note: we may want to
 optimize out this distribution, but for simplicity, I will make this
 directly accessible.

Question: given result $\vec{r} \sim P'(\vec{r}, \vec{i})$ and known \vec{i} is it drawn from $P(\vec{r}, \vec{i})$ or $P_{err}(\vec{r}, \vec{i})$?

if:
 $P(\vec{r}, \vec{i}) > P_{err}(\vec{r}, \vec{i}) \rightarrow P$
 otherwise
 $\rightarrow P_{err}$

Now suppose we do not have \vec{i} . We find which distribution has the highest probability of producing \vec{r} across all inputs.

let $P(\vec{r}) = \int_{\vec{i}} P(\vec{r}, \vec{i}) P(\vec{i})$
 let $P_{err}(\vec{r}) = \int_{\vec{i}} P_{err}(\vec{r}, \vec{i}) P(\vec{i})$
 if
 $P(\vec{r}) > P_{err}(\vec{r}) \rightarrow P(\vec{r}, \vec{i})$
 otherwise
 $\rightarrow P_{err}$

Aside:

We don't have the actual probability distributions: P, P_{err} , we have incomplete approximations \hat{P}, \hat{P}_{err} drawn from a series of observations $(r_{corr}, r_{err}) \in O$.

* We can pay to figure out which distribution the result comes from and improve estimates \hat{P}, \hat{P}_{err} .

1.4 Cost Function

Accuracy Requirement

$$P(|f(\vec{i}) - f'(\vec{i})| < \epsilon) > \omega$$

ϵ : tolerated error

ω : tolerated frequency of intolerable errors.

Relevant machine parameters:

$$P(|f(\vec{i}) - f'(\vec{i})| < \epsilon) \approx p_{err} \int_0^\epsilon P_\epsilon(\epsilon) d\epsilon$$

Compute the probability of an error occuring times the probability the error is less than epsilon.

Semantics for cost function

To differentiate between machines $m \in M : \{m_1, \dots, m_n\} = M$

$$m_i \rightarrow \{P_\epsilon^i, p_{fail}^i, p_{hang}^i, p_{err}^i, c^i\}$$

To differentiate between tasks $t \in T: \{t_1, \dots, t_n\} = T$

There are no particular parameters to annotate, we assume we're talking about one taskset.

The reliable machine has the following properties:

$$r \rightarrow \{P_\epsilon^r = 0, p_{fail}^r = 0, p_{hang}^r = 0, p_{err}^r = 0\}$$

Accuracy Requirement for a task-machine mapping

Given we have assignments, $t_i \in T | t_i \rightarrow m_j$:

$$P(|f(\vec{i}) - f'(\vec{i})| < \epsilon) = \frac{1}{|T|} \sum_{t_i=t_0}^{t_n} P_{m_j}(|f(\vec{i}) - f'(\vec{i})| < \epsilon)$$

The probability of a set of task-machine assignments is the average of each probability of each assignment meeting the assignment.

Minimization formula

For a set of assignments $t_i \in T | t_i \rightarrow m_j$, we want to maximize:

$$\sum_{t_i \in T} c_j$$

Putting it together:

We want to find a set of $t_i \rightarrow m_j$ mappings that minimizes:

$$\sum_{t_i \in T} c_j$$

but meets the following constraint:

$$[\frac{1}{|T|} \sum_{t_i=t_0}^{t_n} P_{m_j}(|f(\vec{i}) - f'(\vec{i})| < \epsilon)] > \omega$$

ie:

$$[\frac{1}{|T|} \sum_{t_i=t_0}^{t_n} p_{err}^j \int_{-\epsilon}^\epsilon P_\epsilon^j(\epsilon) d\epsilon] > \omega$$

#Simple case 1: Unreliable machine, reliable machine, T tasks, no re-execution

Assume, $M = m_0$ and $c_0 < 1, c_r = 1$. The minimization problem becomes run as many tasks on the unreliable machine.

$$[\frac{1}{|T|} \sum_{t_i=t_0}^{t_n} p_{err}^j \int_{-\epsilon}^{\epsilon} P_{\epsilon}^j(\epsilon) d\epsilon] > \omega$$

We can flip the problem to consider the probability of a bad error occurring:

$$[\frac{1}{|T|} \sum_{t_i=t_0}^{t_n} p_{err}^j (1 - \int_{-\epsilon}^{\epsilon} P_{\epsilon}^j(\epsilon) d\epsilon)] < 1 - \omega$$

lets say we execute k tasks on the unreliable processor:

$$[\frac{k}{|T|} p_{err}^j (1 - \int_{-\epsilon}^{\epsilon} P_{\epsilon}^j(\epsilon) d\epsilon)] < 1 - \omega$$

$$k < \frac{|T|(1-\omega)}{p_{err}^j (1 - \int_{-\epsilon}^{\epsilon} P_{\epsilon}^j(\epsilon) d\epsilon)}$$

We can run a maximum of k tasks on the unreliable machine where k is

$$k < \frac{|T|(1-\omega)}{p_{err}^j (1 - \int_{-\epsilon}^{\epsilon} P_{\epsilon}^j(\epsilon) d\epsilon)}$$

1.5 Outlier Detector

Suppose we have an outlier detector $P_{out}(\vec{r})$:

if $P_{out}(\vec{r}) \neq 0 \rightarrow ok$

else $\rightarrow outlier$

Outlier Detector

Suppose the outlier detector perfectly captures the reliable result probability distribution:

$$P_{out}(\vec{r}) = P(\vec{r}) = \int_i P(\vec{r}, \vec{i}) P(\vec{i}) di$$

Impact on Error

The maximum error is bounded by

$$\epsilon_{max} = \operatorname{argmax}_{\vec{r}_a, \vec{r}_b} \sqrt{(\vec{r}_a - \vec{r}_b)^2} \quad | \quad P_{out}(\vec{r}_a) > 0 \text{ and } P(\vec{r}_b) > 0$$

Let $f_e(\epsilon) = \{1 \text{ when } |\epsilon| \leq \epsilon_{max}, 0 \text{ otherwise}\}.$

The new error distribution with this outlier detector is:

$$P_e^{j'}(\epsilon) = P_e^j(\epsilon) \cdot f_e^j(\epsilon)$$

Impact on Cost

The probability of re-execution is

let $g_e(\vec{r}) = \{0 \text{ if } P_{out}(\vec{r}) > 0, 1 \text{ otherwise}\}$

$$p_{reexec}^j = \int_{-\infty}^{\infty} P_{j'}(\vec{r}) \cdot g_e(\vec{r})$$

So the added predicted cost for machine j is

$$c_r \cdot p_{reexec}^j + c_j$$

1.6 Martin Notes

=====

Add probability outlier detector detects outlier (increase cost probabilistically, improve accuracy probabilistically)

=====

Martin:

* 2 different probability distributions, each output comes from a mixture of the two. (statistical inference, figure out likelihood) error

prob distribution, and non-error distribution
 optimization problem: pay cost to get truth.

* joint probability distribution of the inputs and the outputs (deterministic). Project the dots down onto the outputs (integration), figure out which distribution the output comes from.

2 Model

Machine Model

Worker Machine $m \in M, |M|$

Main machine R

Each worker machine has cost:

$$c = a \cdot work \cdot t + b \cdot energy \cdot t$$

energy: energy per time.

work: amount of work per time $t = t_a, t_b$.

For each task T , the machine has the following properties:

For (m, T)

P_{crash} : The per-task runtime probability distribution.

P_{hang} : The per-task runtime probability distribution.

$P_{r_i}(x|r, t)$: The probability distribution of results produced by this machine, given the correct output r and the runtime.

Note: These properties are not known by topaz

Programming Model

For taskset T : $t \in T, |T|$

Task inputs and outputs: $r : R^n \rightarrow R^m \mid t(i_1, i_2 \dots) \rightarrow (r_1, r_2, \dots)$

The inputs are real values. A fixed array of size n is interpreted as n unit values.

each task has a runtime: t

Observed Properties

For each machine:

For $(m, T) \mid m \in M \rightarrow (\{OD_{r_i}\}, OD_{time})$

each task has a time that it takes to execute on some machine $\hat{P}_{r_i}(x)$: For each result r_i , the observed output probability distribution. Used to determine if a runtime is reasonable.

\hat{P}_{time} : The observed distribution of runtimes. Used to determine if a runtime is reasonable.

Scheduling Tasks

$t_1, t_2, \dots \rightarrow m$

topaz may send any number of tasks at once to machine m .

3 Programming Model

4 Notation

Tasks: $t_i \in T, |T| = n$

Machine: $m_j \in M, |M| = m, r : \text{reliable}$

Cost: $m_j^w \in M \mid c_j = a \cdot \frac{1}{\text{speed}_j} + b \cdot \text{energy}_j$

Martin: each machine has a speed and energy. Each task has an amount of work. The amount of time it takes machine to execute task i is $\text{work}_i \cdot \text{speed}_j$.

Topaz executes on a distributed system with m worker machines, $m_i^w \in M, |M| = m$ and a reliable main machine m^r . Each worker machine has a cost associated with it c_i , the reliable machine has a cost c_r .

Optimization Problem: we want to minimize the cost of the execution of the taskset, C , while meeting the quality of result, Q .

4.1 Cost

The cost of a machine $c_i = f(\text{perf}, \text{energy})$, where $c_i \propto \frac{1}{\text{perf}}$ and $c_i \propto \text{energy}$. We can express $c_i = a \cdot \frac{1}{\text{perf}_i} + b \cdot \text{energy}_i$, where the weights may be user specified. I separate the cost from properties of the machine to keep the cost model sufficiently abstract.

Assumption: *The reliable machine is more expensive than any of the other machines in the cluster $\forall c_i : m_i \in M; c_r > c_i$*

4.2 Correctness

We need some way of quantifying the quality of the result. Given a particular solution, each machine generates a result from a distribution centered around the solution. For example, the reliable machine generates the solution with $P(\text{Sol}) = 1$. For now on s_i is the correct solution for task i .

4.2.1 Data Parallel Problems

for problems where each task generates an independent piece of data (ex/pixels in an image), we wish to enforce the following constraint:

Strong Assertion: All tasks meet a result constraint.

$$\forall t_i \in T, t_i \rightarrow m_j \mid (\int_{-\epsilon}^{\epsilon} P_j(s_i + \epsilon)) > P_{spec}$$

ex: black-scholes, we want to ensure each price is likely to give a correct answer

Weak Assertion: On average, all tasks meet a result constraint.

$$t_i \rightarrow m_j \mid \frac{1}{n} \sum_{i=1}^n (\int_{-\epsilon}^{\epsilon} P_j(s_i + \epsilon)) > P_{spec}$$

ex: image generation, allow individual pixels to be arbitrarily bad as long as entire image is on average, ok.

Issues: we don't know the error distribution relative to the solution beforehand, we would need to build a per-machine error distribution during runtime. We can empirically estimate $1 - \int_{-\epsilon}^{\epsilon} P_j(s_i + \epsilon)$ at runtime by re-executing tasks at runtime and taking the number of incorrect results over the total number of results. We will call this p_j^f , the probability of task failure for machine j.

4.3 Assumptions

Martin: No batching then batching. Bad machine generates fault result $U(-inf, inf)$, actual result is $N(0,1)$

- *No Batching:* one task execution on one machine, followed by one reduction.
- *No Catastrophic Failures:* Typically if you send too many tasks to a machine, it may crash. We won't model this for now.
- *Per-Item Outlier Detection:* Simple per-primitive outlier detector - no cross-item correlation.
- *User provides reliability specification:* user provides P_{spec} and ϵ for each data element in each task set.

4.4 Parameter Space

Topaz has the following parameters to play with.

Scheduling: [cost: c_j] We can choose $j : t_i \rightarrow m_j, m_r$:

Outlier Detection: We can configure different parameters of the online outlier detector. For our basic outlier detector, that is the number of standard deviation we can use.

5 Naive Model

More Assumptions:

- *We're discussing a taskset with a single numerical output.*
- *Output Distribution follows Normal Distribution:* A simplification.
- *Data Parallel, Strong Specification:* provide a hard desired probability p_{spec} falls within ϵ

5.0.1 Scheduling

We greedily choose the smallest c_j where $1 - p_j^f$ (the probability of a particular element being outside of $s_i \pm \epsilon$) meets p_{spec}

5.0.2 Outlier Detector

Online detector that models the output distribution as $N(\mu, \sigma)$. This does not work well for output distributions that do not resemble normal distributions. If an outlier detector detects an outlier, for now, always re-execute.

Re-execution: $[\text{cost}:c_j + c_r]$ We can get a completely correct result and improve the outlier detector on outlier detection.

Parameters:

- $K \rightarrow$ the number of standard deviations from mean to consider ok.

Martin analysis: lets take a particular computation where we add up all the numbers first. Outlier detector algorithm (normal distribution), Prove some bound on how results are likely to be. Martin: we will eventually need to include task work - lets leave task work out. Issues we will eventually have to deal with.

- Amount of work in each task (ie tasks have variable work quantities)
- Batching
- Crashes - outlier detection on task times
- 2 kinds of re-executions: pitch result and re-execute. execute someplace else.

Outlier detection strategies: 1. main processor can send outlier data to child machine to remote machine.

6 Example: Array Summation

Assume

- we \sum the results of T tasks, where $t : i \rightarrow r$ and $r \in N(\mu, \sigma)$.
- we have an outlier detector $N(\mu, k\sigma)$ that predicts the output distribution and considers results within k standard deviations ok, where k is unknown.
- on outlier detect, we re-execute
- we have a desired percent error of p
- assume w percent of tasks fail, Assume the output distribution is uniformly distributed $U(F_{min}, F_{max})$.

$$p = \frac{S_{urel} - S_{rel}}{S_{rel}} = \frac{\sum r_{urel} - r_{rel}}{\sum r_{rel}}$$

All undetected outliers fall within $(\mu - k\sigma, \mu + k\sigma)$. In the worst case: $N(\mu, k\sigma) \rightarrow |r_{urel} - r_{rel}| \sim 2k\sigma$

With increasing k , the percent of undetected outliers increases:

$$\int_{\mu-k\sigma}^{\mu+k\sigma} \frac{1}{F_{max}-F_{min}} \\ [(\mu + k\sigma - \mu + k\sigma) \frac{1}{F_{max}-F_{min}}] \\ [\frac{2k\sigma}{F_{max}-F_{min}}]$$

We can approximate the absolute error as follows:

$$\sum r_{urel} - r_{rel} \approx w \cdot |T| \cdot \frac{2k\sigma}{F_{max}-F_{min}} \cdot k\sigma = w \cdot |T| \cdot \frac{2k^2\sigma^2}{F_{max}-F_{min}}$$

We can approximate the average sum as follows:

$$\sum r_{rel} = |T|\mu$$

We rewrite p as:

$$p \approx \frac{w \cdot |T| \cdot 4k^2\sigma^2}{(F_{max}-F_{min}) \cdot |T|\mu} \\ p \approx \frac{w \cdot 4k^2\sigma^2}{(F_{max}-F_{min})\mu} \\ \frac{(F_{max}-F_{min})\mu \cdot p}{w \cdot 2\sigma^2} = k^2$$

Therefore, the ideal number of standard deviations to consider, given a desired output error rate, is:

$$k = \sqrt{\frac{(F_{max}-F_{min})\mu \cdot p}{w \cdot 4\sigma^2}}$$

Obviously (1) we don't know the error distribution, nor does (2) our outlier detector perfectly fit the output distribution. The the rate of task errors for a particular machine is not known.